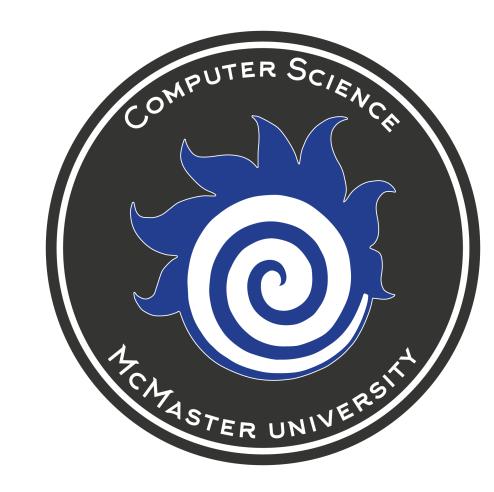# Hakaru Language: Standard Library Implementation and Language Validation Testing

Justin Staples, Mahmoud Khattab, Nevin Mahilal and Aryan Sohrabi, supervised by Dr. Christopher Anand & Dr. Jacques Carette

{staplejw, khattm, mahilank, sohraa3, anandc, carette}@mcmaster.ca

Department of Computing and Software, McMaster University

Introduction

Often, we wish to create a model for some kind of real world phenomenon so that we can extract useful information and learn more about it. For example, in a machine learning application, the goal is often to predict or infer information about the model based on the past outcomes of some kind of experiment. An experiment could result in many different outcomes, each one with a different likelihood (probability).

The mathematical functions that we use to describe these probabilities are called probability distributions and the quantity that denotes the outcome of the experiment is called a random variable. A random variable can take on continuous or discrete values depending on the application. Continuous random variables are formally defined by a probability density function (PDF), which maps the outcomes of the random variable to real numbers that represent a probability. Similarly, a discrete random variable is described by a probability mass function (PMF). As an example, let us consider one of the most ubiquitous and naturally occuring distributions, the normal distribution. We can say that a random variable, X, is normally distributed with the following notation.

$$X \sim Normal(\mu, \sigma^2)$$

Here, we can see that the normal distribution is parametrized by μ, the population mean, and σ, the standard deviation. A plot of the PDF is shown below. Remember that the PDF describes the liklihood that a value sampled from this distribution will equal x.
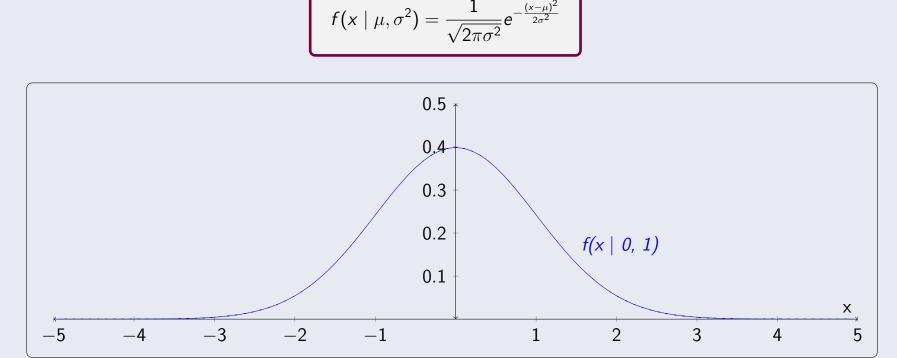
$$f(x \mid \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Figure 1: PDF of the normal distribution, with $\mu = 0$ and $\sigma = 1$

Hakaru is an experimental **probabilistic programming language**, which aims to simplify the way users implement statistical distributions.

- Niche application means that Hakaru is a small language with limited features.
- Running a Hakaru program generates a stream of random numbers distributed according the statistical distribution modeled.
- Hakaru code can be compiled to C and Haskell so that models can be exported for use in larger applications.
- Terminology: implementation of statistical distribution $\iff$ probabilistic model $\iff$ model.

'hk-maple' is a provided inference algorithm which takes a hakaru program file as an argument and uses Maple to perform algebraic transformations. If the Simplify mode of 'hk-maple' is used (the default mode), it returns an equivalent hakaru program with greater sampling efficiency. In applications requiring billions of samples (e.g. machine learning) this has the potential to save significant processing time!

Key Concepts

This section is dedicated to explaining some key ideas and terminology for our project.

- ⋆ Hakaru program $\iff$ implementation of statistical distribution $\iff$ probabalistic model $\iff$ model.
- ⋆ Values pulled from a distribution are called measures. Measures $\iff$ samples.

- ⋆ '∼' vs '<∼':
  - $X \sim Normal(0, 1)$ means that the random variable, X, is distributed according to the standard normal distribution (statistics literature).
  - x <∼ normal(0, 1) means to pull a random variable from the standard normal distribution and bind it to the variable, x (Hakaru code).
- ⋆ PDF: Probability Density Function (continuous random variables).
- ⋆ PMF: Probability Mass Function (discrete random variables).
- ⋆ UDR chart: Univariate Distribution Relationship chart (describes relationships and transformations between different distributions).

Motivation

OBJECTIVES:

- ⋆ Increase language accessibility: One of our main aims was to add more features to the language such as primivite mathematical functions (e.g. log, choose), new distributions (standard library), error handling, etc.
- ⋆ Test language validity: The results of the 'hk-maple' tests that we run can validate the Hakaru language if the transformations are recognized by the simplifier.

ENDEAVOURS:

- ⋆ Standard library development: This mainly involves using the UDR to branch out and implement new distributions.
- ⋆ Syntax-highlighting-for-hakaru Package for Sublime Text: This is a package developed by Mahmoud. It makes developing Hakaru code so much more convenient and streamlined.
- ⋆ Test case writing: There are a plethora of test cases to consider with these distributions, so we only focus on exhausting the list for a select few (e.g. chi-square).

Standard Library Development

In order to make Hakaru more accessible to users, the standard library should implement probabilistic models of statistical distributions that are commonly used by statisticians. To this end, we have implemented the distributions described in the UDR, a portion of which can be seen below. In general, we have used the following principles to guide the development of the standard library:

- ⋆ Whenever possible, implement distributions as transformations on pre-existing models.
- ⋆ Reference the UDR chart for possible ways to implement each distribution.
- ⋆ In the case of multiple possible implementations, take the shortest possible path from a primitive distribution on the UDR.
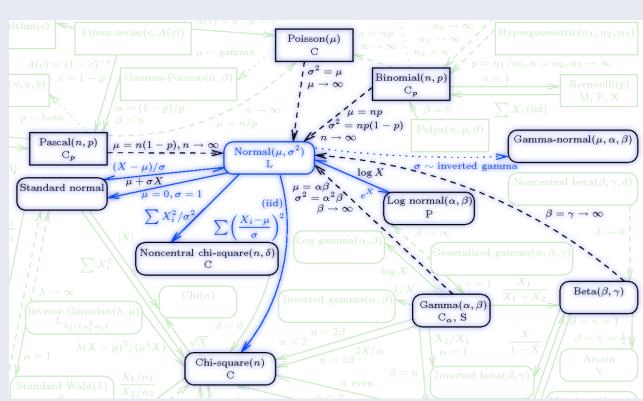

Figure 2: A snapshot of the UDR shows how the normal distribution can be transformed into a multitude of other distributions.

Hakaru functions representing probabilistic models (i.e. functions with a measure(<type>) return type) cannot be used as operands in normal mathematical operations. This is to say we cannot transform models directly (i.e. we cant do algebraic transformation on the PDF). The only operator we can use on a model is bind (<∼), to pull a sample from it. However, we are able to transform samples however we like. Therefore, we are interested in implementing transformations of the following form.

$$R(p, q) \Rightarrow X \sim A(p) \Rightarrow f(X) \sim B(q)$$

In the equation above, p is a set of values parameterizing the distribution A, q is a set of values parameterizing the distribution B, $R(p, q)$ is a set of relationships between p and q that must be satisfied, and f is a function that applies a transformation to a sample. We can expand this definition to include transformations defined in terms of an aggregation of multiple independent samples. For example, the standard chi-square distribution is defined as the sum of the squares of n normal random variables (see Figure 3).

```
def chiSq(means array(real), stdevs array(prob) ):
    q <~ plate _ of size(means): normal(means[_],stdevs[_])
    return summate i from 0 to size(q):
        ((q[i]-means[i])/stdevs[i])^2
```

Figure 3: Our implementation of the chi-square distribution.

Hakaru also lends itself very well to Bayesian transformations where we want to pull a sample from a distribution which is parameterized with a value sampled from another distribution. These transformations take the following form. Here, we are saying that if X is parameterized by p, then the distribution C, which is parameterized by q and X, is equal to another distribution, B, which is parameterized by q and X. The gamma-poisson distribution can be described by such a transformation (see Figure 4).

$$X \sim A(p) \Rightarrow Y \sim B(q, X) = C(p, q)$$

```
def gammaPoisson(shape prob, scale prob) measure(nat):
    mu <~ gamma(shape, scale)
    X <~ poisson(mu)
    return X
```

Figure 4: Our implementation of the gamma-poisson distribution.

Some distributions on the UDR are unreachable from Hakaru's primitive distributions using transformations like the 2 described above. In these cases, we have to implement the model in terms of the PMF for a discrete distribution or in terms of the PDF for a continuous distribution. We now discuss this process in the following examples.

Discrete Distributions from PMF

When the PMF implies a finite number of categories, the use of the primitive categorical distribution combined with the array constructor language feature lends itself very well to an implementation of the distribution. However, discrete distributions have PMFs implying an infinite number of categories. This cant be implemented using categorical, which requires a finite size array. In this case, we use the weight function primitive to Hakaru. The weight function assigns a weight to each sample. Shown below is an example using the weight function.

```
burglary <~ categorical([0.0001, 0.9999])
weight([0.95, 0.01][burglary],
return [true,false][burglary])
```

In this example, the value burglary is sampled from a categorical distribution. Then, the samples are reweighted so that true has a weight of 0.95 and false has a weight of 0.01. For the distribution samplings that are implemented in the standard library, the first argument is the PMF or PDF of the distribution and the second argument is the random variable. For example the following is the code that implements a model of the discrete Weibull distribution

```
def discreteWeibull(p prob, beta prob):
    x <~ counting()
    weight(if (x < 0): 0 else: (1 - p) ** (x ** beta)
              - (1 - p) ** ((x + 1) ** beta), return x)
```

The first line of code in the function definition, x <∼ counting(), assigns an integer to x according to a uniform distribution across all integers. The first argument of weight is the PMF of the distribution and the second argument is the random variable. In this way, each instance of the random variable is given a weight equal to that instance's probability.

Continuous Distributions from PMF

This is accomplished similar to how discrete distributions with infinite categories implied by their PMF are implemented. The main difference is that the lebesgue function is used instead of counting to sample values from the real number line.

Testing Relationships Between Distributions

Recall that when developing the standard library, the UDR chart often implied multiple possible implementations for a given distribution. Naturally, we would expect alternative implementations for the same distribution to result in equivalent probabilistic models. This line of thought is the basis for the kinds of test cases we have implemented. More generally, we have the following hypothesis.

Assume we know a proven relationship between 2 statistical distributions, A and B, which allows us to transform A and B into distributions that are equivalent to each other. Further assume this transformation takes on one of the forms discussed in the Standard Library Development section. We hypothesize that by applying the appropriate transformations to implementations of A and B, we can create two Hakaru programs whose 'hk-maple' outputs will be equivalent to each other. Test cases that prove our hypothesis true indicate the validity of the Hakaru language implementation. Test cases that prove our hypothesis false indicate an underlying bug in the language definition which is to be passed back to the language developers.

Conclusions & Future Work